



IMAGE CLASSIFICATION USING CNN AND TRANSFER LEARNING

By: Anirudh Unni, Natalia Dominguez and Zahory Vasquez

Overview

| | |
|-------------------------------|----|
| ▶▶▶ Project Overview | 03 |
| ▶▶▶ Dataset Description | 04 |
| ▶▶▶ Data Preprocessing | 05 |
| ▶▶▶ Custom CNN | 04 |
| ▶▶ Architecture & Training | |
| ▶▶ Evaluation & Improvements | |
| ▶▶ CNN Evaluation | 05 |
| ▶▶ Transfer Learning Approach | 06 |
| ▶▶ Performance Comparison | 07 |
| ▶▶ Conclusions | 08 |
| ▶▶ Future Work | 09 |



PROJECT OVERVIEW

Goal:

- Classify images from the CIFAR-10 dataset into 10 categories.

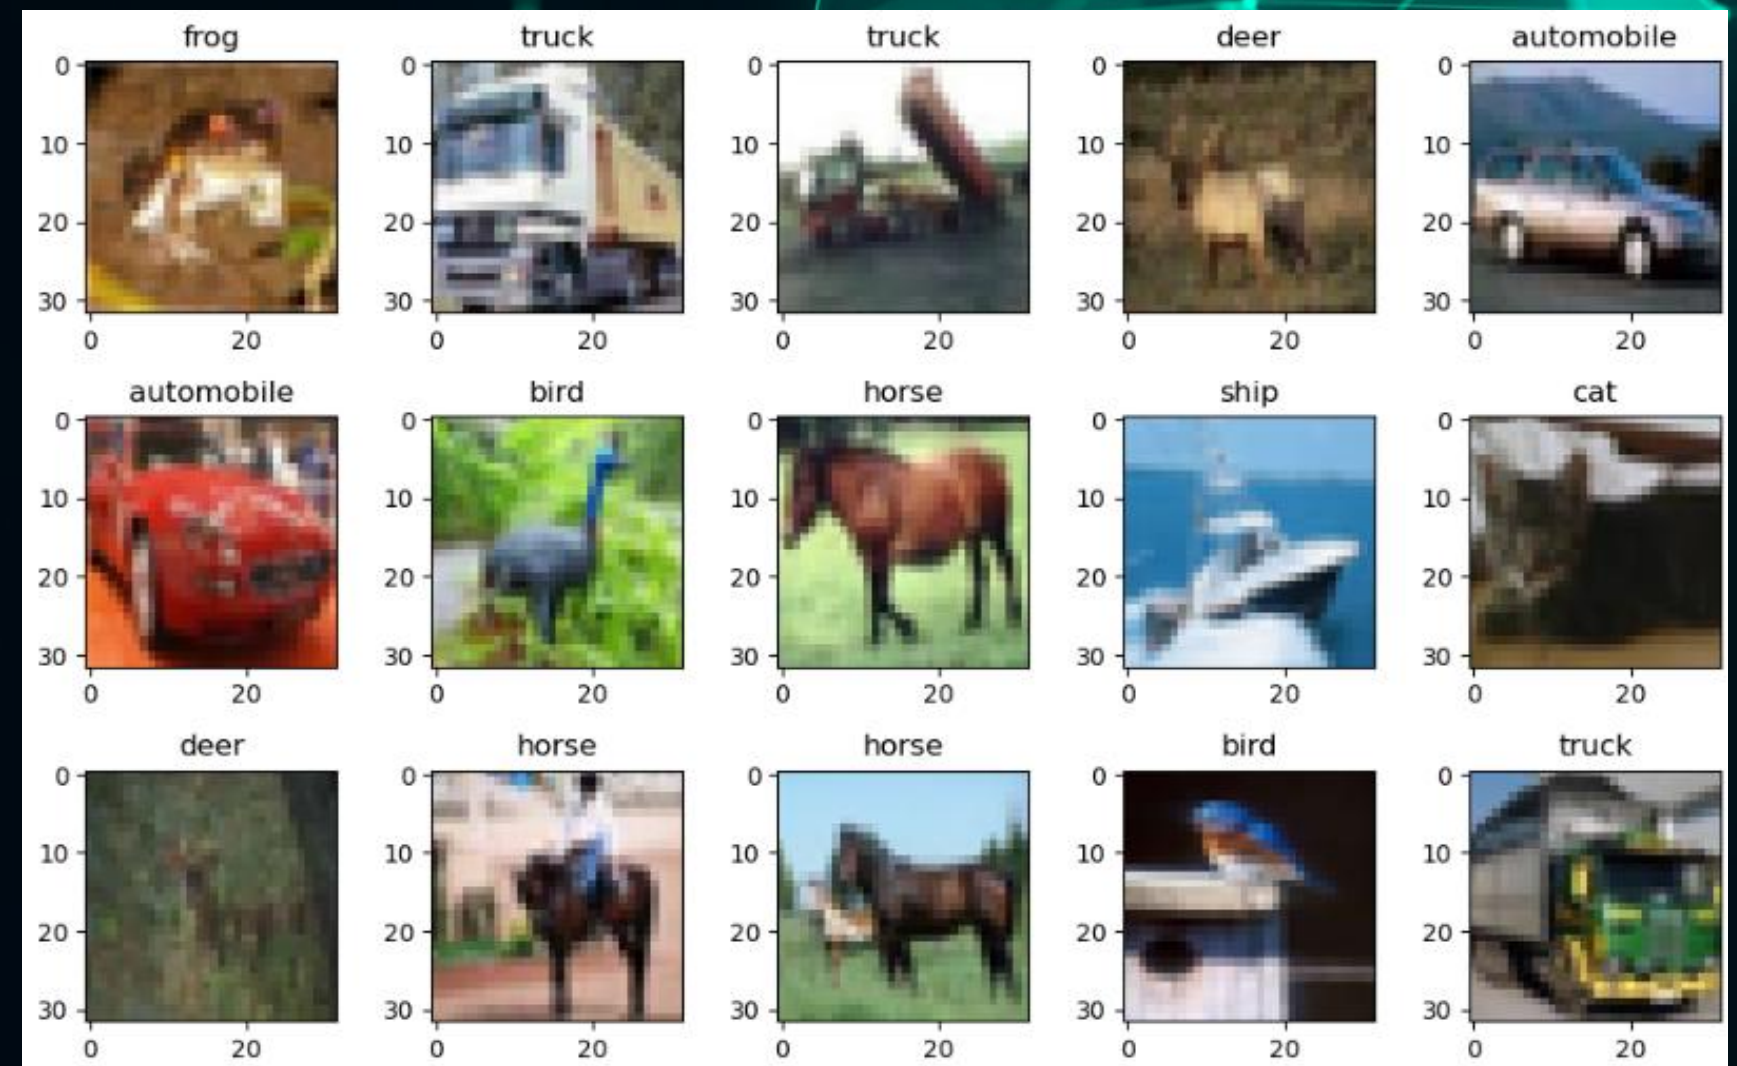

Approaches:

- Custom CNN built from scratch.
- Transfer Learning using pre-trained models.



DATASET DESCRIPTION

- **Dataset:** CIFAR-10
- **Size:** 60,000 color images (32×32 pixels).
- **Classes:** 10 categories



DATA PREPROCESSING

Before training, the dataset must be prepared to fit the model:

1. **Reshape**: Adjust input images to include the channel dimension required by CNNs.
2. **Normalize**: Scale pixel values to the $[0, 1]$ range to improve training efficiency.
3. **Encode**: Convert class labels into one-hot vectors for multi-class classification.

Implementation uses [Keras](#) utilities:

- Rescaling for normalization
- `to_categorical` for label encoding



CUSTOM CNN ARCHITECTURE & TRAINING

Model Architecture:

- Convolutional layers with ReLU activations
- MaxPooling layers for feature reduction
- Flatten layer → Fully Connected dense layers

Model Training:

- Optimizers: Adam
- Loss function: Categorical Crossentropy

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 30, 30, 32) | 896 |
| max_pooling2d (MaxPooling2D) | (None, 15, 15, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 13, 13, 64) | 18,496 |
| max_pooling2d_1 (MaxPooling2D) | (None, 6, 6, 64) | 0 |
| flatten (Flatten) | (None, 2304) | 0 |
| dense (Dense) | (None, 128) | 295,040 |
| dense_1 (Dense) | (None, 10) | 1,290 |

Total params: 315,722 (1.20 MB)

Trainable params: 315,722 (1.20 MB)

Non-trainable params: 0 (0.00 B)

CUSTOM CNN

EVALUATION & IMPROVEMENTS


Model Evaluation:

- **Metrics:** Accuracy, Precision, Recall, F1-score
- Confusion matrix to analyze class-level performance

Test

Accuracy: 0.6748
Precision: 0.6841
Recall: 0.6748
F1 Score: 0.6761

Training accuracy: 0.82



CUSTOM CNN

EVALUATION & IMPROVEMENTS

Improving the Model:

- Added more convolutional layers for deeper feature extraction
 - Added 1 more layer (previously only had 2) + Dropout

Test accuracy with more layers: 0.7343

- Added BatchNormalization + GlobalAveragePooling2D

Test accuracy with more layers: 0.7313

- Added Regularizer

Test accuracy with more layers: 0.7249

- Added Hyperparameter turning - Earling Stopping

Test accuracy with early stopping: 0.7293

- Applied data augmentation (rotations, flips, shifts) to improve generalization

Test accuracy with data augmentation: 0.7694

Training accuracy: 0.85

Training accuracy: 0.91

Training accuracy: 0.93

Training accuracy: 0.94

Training accuracy: 0.74

TRANSFER LEARNING APPROACH

Instead of evaluating multiple architectures, our approach focused exclusively on MobileNetV2 (pre-trained on ImageNet). The model was tested in three different configurations, varying which layers were frozen and the extent of fine-tuning, in order to compare performance and training time.



METRICS & EVALUATION

Key Features

Results

MobileNetV2.1

- **Simple dense head:** GlobalAveragePooling → Dense(128, ReLU) → Dense(10, softmax)
- **Training:** 5 epochs, batch size 64

Test accuracy: 0.8658

MobileNetV2.2

- **Dense head with dropout:** GlobalAveragePooling → Dropout(0.3) → Dense(10, softmax)
- **Training:** Frozen base first phase, 10 epochs, default batch size

Test accuracy: 0.8562

MobileNetV2.3

- Same architecture as V2.2
- **Training:** Frozen base first phase, 10 epochs, batch size 64

Test accuracy: 0.8557

CONCLUSIONS

Custom CNN (with Data Augmentation):

- Solid baseline model.
- Achieved moderate performance with accuracy = 0.7694.
- Requires longer training time since all weights are learned from scratch.

Transfer Learning (MobileNetV2.1):

- Achieved higher accuracy = 0.8658.
- Faster convergence by leveraging pre-trained features.
- Better generalization compared to training from scratch.

Takeaway:

- Transfer learning is the recommended approach for similar image classification tasks.

FUTURE STEPS

- Test other pre-trained models (e.g., EfficientNet, ResNet).
- Fine-tune more layers for better feature learning.
- Apply advanced data augmentation techniques.
- Optimize hyperparameters for higher accuracy.

